

The University of Manchester  
Department of Computer Science  
Project Report 2024

**My Little Operating System**

Author: Thomas H. Jones

Supervisor: Dr James Garside

**Abstract**

My Little Operating System

Author: Thomas H. Jones

Supervisor: Dr James Garside

# Contents

<b>1</b>	<b>Introduction</b>	<b>6</b>
1.1	Aims . . . . .	6
1.2	Motivations . . . . .	6
1.3	Project Plan . . . . .	7
1.4	Report Outline . . . . .	7
<b>2</b>	<b>Background</b>	<b>8</b>
2.1	Operating System Fundamentals . . . . .	8
2.1.1	Process Management . . . . .	8
2.1.2	Memory Management . . . . .	9
2.1.3	IO Management . . . . .	10
2.2	Information on board . . . . .	10
2.3	Risc-v fundamentals . . . . .	10
2.3.1	Extensions . . . . .	10
2.3.2	Traps . . . . .	10
2.3.3	Important CSRs . . . . .	10
2.3.4	Comparison between RISC-V and ARM . . . . .	10
<b>3</b>	<b>Design and Implementation</b>	<b>11</b>
3.1	Toolchain . . . . .	11
3.2	Board Configuration . . . . .	11
3.2.1	Clock settings . . . . .	11
3.3	Processes . . . . .	13
3.3.1	Process Structure . . . . .	13
3.3.2	Scheduler . . . . .	14
3.4	Memory Management . . . . .	14
3.5	IO . . . . .	14

<b>4 Outcome</b>	<b>15</b>
<b>5 Evaluation</b>	<b>16</b>
<b>6 Conclusions and further work</b>	<b>17</b>
<b>Bibliography</b>	<b>18</b>

# List of Figures

3.1	High Frequency Clock Diagram . . . . .	12
3.2	Low Frequency Clock Diagram . . . . .	12

## List of Tables

# Abbreviations

**CLINT** core local interruptor

**IO** Input/Output

**ISA** instruction set architecture

**PLIC** platform level interrupt controller

# Chapter 1

## Introduction

### 1.1 Aims

This project will be based upon the following goals

- A basic implementation of processes, which can run arbitrary user code
- A scheduler that schedules processes interactively
- Memory management, to ensure that user code has limited access to memory
- Input/Output (IO), for example a text input and output, or other visual methods of output

### 1.2 Motivations

This project is an exploration of concepts and methods of operating systems, with the overall goal to implement these concepts to create a simple operating system. In addition to this, the project will be done using a board that implements the open source, RISC-V processor architecture, which will require research into the operation of RISC-V, as well as any implementation specific features to the board in use. The board used will be the Sifive HiFive1 Rev B. The main goal is the implementation of an interactive scheduler to run multiple processes at once, as this is something that will not be implemented unless part of an operating system.



## 1.3 Project Plan

The initial part of this project will be determining an appropriate microcontroller board and setting up the toolchain required to compile and load code onto the board. Following this, the basic components of the system will be set up, such as the interrupt handler. Simple IO will be setup, to allow for ease of debugging, which will then be reimplemented later to be consistent with other IO elements. The implementation of processes will be split into multiple parts, such as the storage of process information, the creation of processes, the scheduling of processes and the deletion of processes. As part of this, memory management will be developed at the same time, as a basic implementation is required for user code to function.

## 1.4 Report Outline

- Introduction: This chapter gives the basic overview of the planned project
- Background: This chapter will introduce key concepts to enable a more comprehensive understanding of the projects details
- Design: This chapter will outline the key design decisions made before and during the project
- Implementation: This chapter will give details on the practical implementation of concepts and designs mentioned in previous chapters
- Evaluation: This chapter will discuss the implementation and findings made during the implementation
- Conclusion: This chapter will give a summary of how the aims of the project were met

# Chapter 2

## Background

This chapter will go over the background of this project, which will cover basic operating system fundamentals, a brief overview of the RISC-V architecture, and the specifics of the board that will be used, which is the Sifive HiFive1 RevB. By the end of the chapter, the reader should be able to understand how parts of the operating system should function, and be able to understand the RISC-V architecture in comparison to the ARM architecture.

### 2.1 Operating System Fundamentals

An operating system's purpose is to provide an interface between user code and the hardware, such as memory and IO, to allow the user code to function seamlessly, and allow interaction with users. This has been split into three key sections, as listed by the aims of the project, as process management, memory management and IO management.[2]

#### 2.1.1 Process Management

The process model is where all software is organised as separate code, which is scheduled in a fashion to give the illusion of multiprocessing. While multiprocessing is possible with multiple hardware threads, this project is limited to microcontroller chips, which mainly only support a single hardware thread. To give the illusion of multiprocessing, the information to load and run individual sections of sequential code is stored, so that an algorithm can be used to load, run, and possibly pause these execution, to allow them to be swapped out at a

rate that makes it imperceivable that they are not running concurrently. This algorithm is the scheduling algorithm, and there are several approaches to how this is handled.[2]

### **Scheduling Algorithms**

In an interactive system, processes are ran in a way that ensures each process will get a fair amount of execution time. This is done simplest using a round robin algorithm. Due to the interactive requirement, processes that take a large amount of execution time need to be pre-empted, so one process cannot halt other processes for too long. This is done mainly by pausing a running process after a certain amount of time (which is known as the algorithms quantum).

#### **2.1.2 Memory Management**

The goal of memory management is to streamline how a process can use memory, while at the same time protecting critical sections of memory from faulty or malicious user code. In a system with multiple processes being executed and no memory abstraction, there exists the possibility that two processes attempt to use the same section of memory, creating a conflict that would cause both processes to run incorrectly. This occurs as different processes cannot be aware of each other, and have no choice but to use memory without knowledge of which sections are in use. This can be solved using memory abstraction. The simplest method of this is using address spaces, where each process is given permission to access only segments of directly addressed memory. This prevents each process modifying other processes memory, and allows the process to behave individually, as long as it is provided with the location of its address space.[2]

### **2.1.3 IO Management**

## **2.2 Information on board**

## **2.3 Risc-v fundamentals**

### **2.3.1 Extensions**

The Hifive1 uses a 32 bit E31 core, which runs the RV32IMAC instruction set architecture (ISA), with user and machine privilege levels.

### **2.3.2 Traps**

In RISC-V, a trap refers to anything where the execution on the hart is handed to the trap handler. There are two categories of traps, synchronous and asynchronous

#### **Asynchronous/Interrupts**

An asynchronous trap is a break in execution caused by external factors, and can be referred to as an interrupt. There are three types of interrupt, software, timer, and external. These are controlled by two units, the core local interruptor (CLINT) and the platform level interrupt controller (PLIC). The CLINT handles the software and timer interrupts, and the PLIC handles the external interrupts.

#### **Synchronous**

A synchronous trap occurs in response to the execution of an instruction, which occurs with the clock, hence synchronous. These come from two sources, errors and environment calls. Errors include memory misalignments, access faults and illegal instructions. This allows for these errors to be handled, either to inform the process of the error, or to remove the process. Environment calls are how processes make calls to the kernel to perform actions that are above the process access level, generally to perform IO operations, or to interact with other processes.

### **2.3.3 Important CSRs**

### **2.3.4 Comparison between RISC-V and ARM**

# **Chapter 3**

## **Design and Implementation**

### **3.1 Toolchain**

### **3.2 Board Configuration**

On the Hifive1, there are several important configuration options that affect general operation of the board. The most notable of these are the clock settings, as these indicate the frequency of the processor, input and output frequencies, and timer interrupts.

#### **3.2.1 Clock settings**

The Hifive1 has 3 clock regions, a high frequency clock, a low frequency clock, and a clock used to drive the JTAG connection. The JTAG driver is constant and only used for debugging through JTAG, so is not relevant here.



the realtime clock and the machine timer, both of which are used to generate timed interrupts. Similar to the high frequency clock can be driven from a ring oscillator or from an external clock, which in the Hifive1 is a crystal oscillator. The low frequency ring oscillator functions at 1.5 KHz to 230 KHz using a frequency divider, and the implemented external clock runs at a constant 32.768 KHz, with no option to divide the frequency.

For both clock domains, the crystal oscillator was chosen. The ring oscillator gives the option to operate at a higher frequency, which would result in a higher number of operations per second. While in a practical operating system this would be desirable, since this system is not intended for practical use, a constant frequency was more desirable as it would give more predictable results, and makes IO operations more reliable. For the low frequency clock, a high frequency would be beneficial for a real time system as the higher frequency would allow for more precise timing of interrupts and other functions, however for an interactive system this precision is not required, and so similar to the high frequency domain the constant frequency of a crystal oscillator was selected.

## **3.3 Processes**

### **3.3.1 Process Structure**

Due to the extremely limited amount of memory available on the Hifive1, the data required to store information on each process has to be structured carefully, else the amount of memory required to store process information would begin to limit the amount of memory available to the processes themselves. In the current implementation, 35 bytes are used to store process information.

One byte each is allocated for a process id, process parent id, process status, and size of a processes text section. For process id and process parent id, only one byte is needed as the Hifive1 does not have the memory to support a large of processes, so a theoretical cap of 256 processes is acceptable. Process status can only take 4 values, so only 2 bits of the byte are used, and the text size byte stores the size as a power of 2, where a process has  $2^n$  words, where  $n$  is the value stored. While some of these values only require bits, they are stored as bytes to keep the structure word aligned.

One word each is allocated for the processes program counter, text section pointer and address space pointer. Since memory addresses are word length, these cannot be reduced.

The vast majority of the process structure is used storing the 31 general purpose registers. This is required to retain the state of each process in between scheduling. The only option to reduce this would be to limit the amount of registers available to use. Only 31 must be stored as the x0 register is hardwired to zero.

In other systems, information like the processes stack pointer may be stored, however standard RISC-V calling convention specifies x2 to be used as the stack pointer, so separate storing of this information is not required, and allows a process to handle its address space on its own, however on process creation x2 and x3 are initialised as the stack pointer and the global pointer, where the stack pointer points to the bottom of the process address space and the global pointer at the top. This is done to reduce the overhead of processes that use the standard calling convention.

### **3.3.2 Scheduler**

To implement an interactive system, a scheduler will be designed to allow a process run for a limited time, before halting the process, and running another ready process. The time allowed for a process to run is referred to as its quantum. Processes will be scheduled to be ran in a round robin fashion, to allow a fair runtime to each process while maintaining a low overhead. This is implemented using a circular queue, where new and halted processes are added to the end of the queue, and processes to run are taken from the start of the queue.

## **3.4 Memory Management**

## **3.5 IO**



## **Chapter 4**

### **Outcome**

# **Chapter 5**

## **Evaluation**

## **Chapter 6**

### **Conclusions and further work**

# Bibliography

- [1] *SiFive FE310-G002 Manual v1p5*. Sifive Inc, 2022.
- [2] H. B. Andrew S. Tanenbaum. *Modern Operating Systems 4th Edition*. Pearson, 2015.